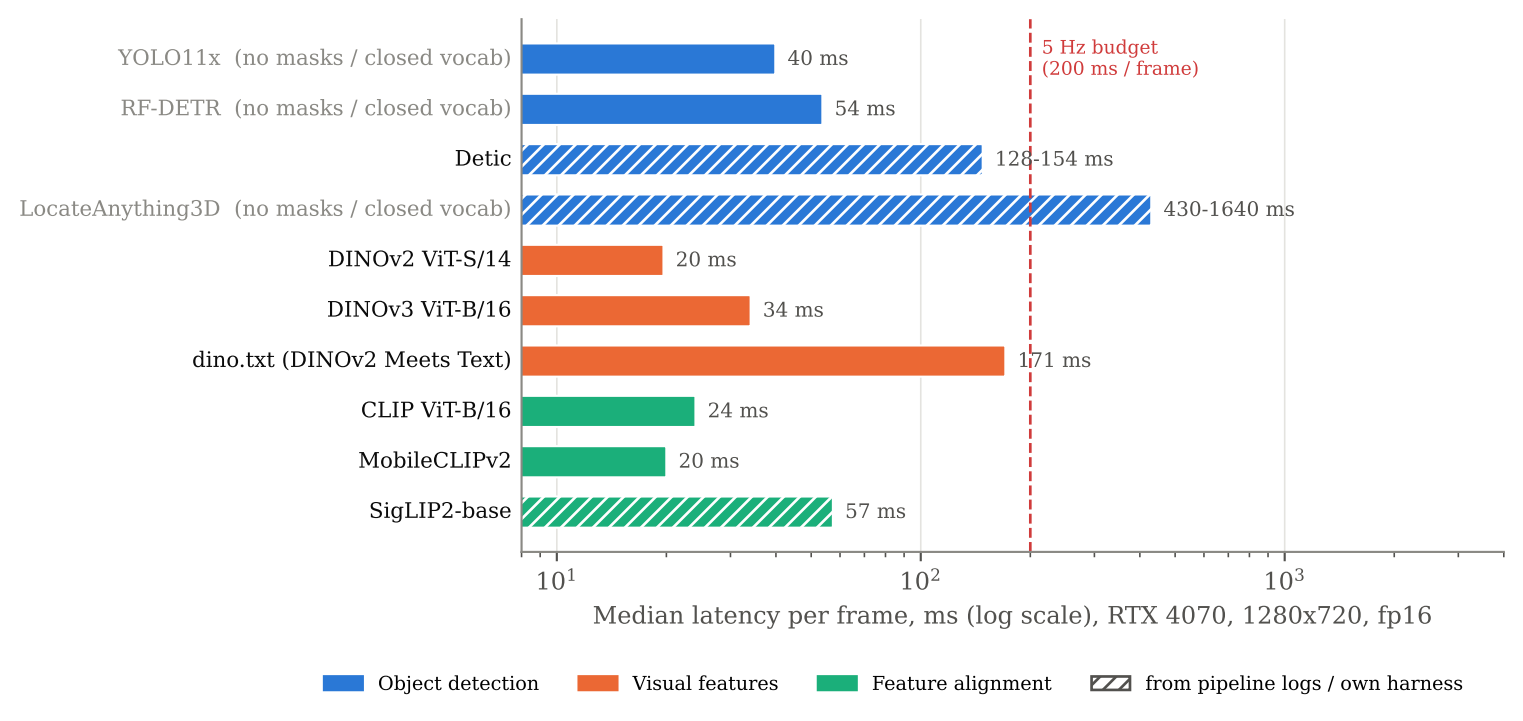
YOLO11x (no masks / closed vocab)
RF-DETR (no masks / closed vocab)
Detic
LocateAnything3D (no masks / closed vocab)
DINOv2 ViT-S/14
DINOv3 ViT-B/16
dino.txt (DINOv2 Meets Text)
CLIP ViT-B/16
MobileCLIPv2
SigLIP2-base
40 ms
54 ms
128-154 ms
430-1640 ms
20 ms
34 ms
171 ms
24 ms
20 ms
57 ms
5 Hz budget
(200 ms / frame)
10^1
10^2
10^3
Median latency per frame, ms (log scale), RTX 4070, 1280x720, fp16
Object detection
Visual features
Feature alignment
from pipeline logs / own harness